

Linear classification

Course of Machine Learning
Master Degree in Computer Science
University of Rome "Tor Vergata"
a.a. 2020-2021

Giorgio Gambosi

Classification

- value t to predict are from a discrete domain, where each value denotes a *class*
- most common case: disjoint classes, each input has to assigned to exactly one class
- input space is partitioned into *decision regions*
- in *linear classification models* decision boundaries are linear functions of input \mathbf{x} ($D - 1$ -dimensional hyperplanes in the D -dimensional feature space)
- datasets such as classes correspond to regions which may be separated by linear decision boundaries are said *linearly separable*

Regression and classification

- Regression: the target variable \mathbf{t} is a vector of reals
- Classification: several ways to represent classes (target variable values)
- Binary classification: a single variable $t \in \{0, 1\}$, where $t = 0$ denotes class C_0 and $t = 1$ denotes class C_1
- $K > 2$ classes: "1 of K " coding. \mathbf{t} is a vector of K bits, such that for each class C_j all bits are 0 except the j -th one (which is 1)

Approaches to classification

Three general approaches to classification

1. find $f : \mathbf{X} \mapsto \{1, \dots, K\}$ (*discriminant function*) which maps each input \mathbf{x} to some class C_i (such that $i = f(\mathbf{x})$)
2. *discriminative approach*: determine the conditional probabilities $p(C_j|\mathbf{x})$ (*inference phase*); use these distributions to assign an input to a class (*decision phase*)
3. *generative approach*: determine the class conditional distributions $p(\mathbf{x}|C_j)$, and the class prior probabilities $p(C_j)$; apply Bayes' formula to derive the class posterior probabilities $p(C_j|\mathbf{x})$; use these distributions to assign an input to a class

Discriminative approaches

- Approaches 1 and 2 are *discriminative*: they tackle the classification problem by deriving from the training set conditions (such as decision boundaries) that, when applied to a point, discriminate each class from the others
- The boundaries between regions are specified by *discrimination functions*

Generalized linear models

- In linear regression, a model predicts the target value; the prediction is made through a linear function $y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$ (linear basis functions could be applied)
- In classification, a model predicts probabilities of classes, that is values in $[0, 1]$; the prediction is made through a *generalized linear model* $y(\mathbf{x}) = f(\mathbf{w}^T \mathbf{x} + w_0)$, where f is a non linear *activation function* with codomain $[0, 1]$
- boundaries correspond to solution of $y(\mathbf{x}) = c$ for some constant c ; this results into $\mathbf{w}^T \mathbf{x} + w_0 = f^{-1}(c)$, that is a linear boundary. The inverse function f^{-1} is said *link function*.

Generative approaches

- Approach 3 is *generative*: it works by defining, from the training set, a *model* of items for each class
- The model is a probability distribution (of features conditioned by the class) and could be used for random generation of new items in the class
- By comparing an item to all models, it is possible to verify the one that best fits

1 Discriminant functions

Linear discriminant functions in binary classification

- Decision boundary: $D - 1$ -dimensional hyperplane $y(\mathbf{x}) = 0$ of all points s.t. $\mathbf{w}^T \mathbf{x} + w_0 = 0$
- Given $\mathbf{x}_1, \mathbf{x}_2$ on the hyperplane, $y(\mathbf{x}_1) = y(\mathbf{x}_2) = 0$. Hence,

$$\mathbf{w}^T(\mathbf{x}_1) - \mathbf{w}^T(\mathbf{x}_2) = \mathbf{w}^T(\mathbf{x}_1 - \mathbf{x}_2) = 0$$

that is, $\mathbf{x}_1 - \mathbf{x}_2, \mathbf{w}$ orthogonal

- For any \mathbf{x} s.t. $y(\mathbf{x}) = 0$, $\mathbf{w}^T \mathbf{x}$ is the length of the projection of \mathbf{x} in the direction of \mathbf{w} (orthogonal to the hyperplane $y(\mathbf{x}) = 0$), in multiples of $\|\mathbf{w}\|_2$
- By normalizing wrt to $\|\mathbf{w}\|_2 = \sqrt{\sum_i w_i^2}$, we get the length of the projection of \mathbf{x} in the direction orthogonal to the hyperplane, assuming $\|\mathbf{w}\|_2 = 1$
- Since $\mathbf{w}^T \mathbf{x} = -w_0$,

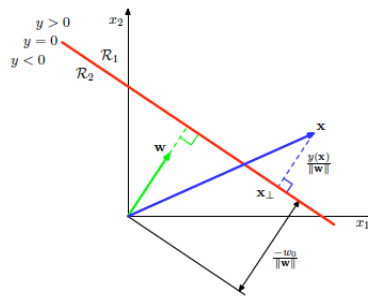
$$\frac{\mathbf{w}^T \mathbf{x}}{\|\mathbf{w}\|} = -\frac{w_0}{\|\mathbf{w}\|}$$

thus, the distance is determined by the threshold w_0

Linear discriminant functions in binary classification

- In general, for any \mathbf{x} , $y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$ returns the distance (in multiples of $\|\mathbf{w}\|$) of \mathbf{x} from the hyperplane

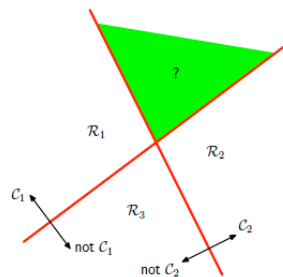
- The sign of the returned value discriminates in which of the regions separated by the hyperplane the point lies



Linear discriminant functions in multiclass classification

First approach

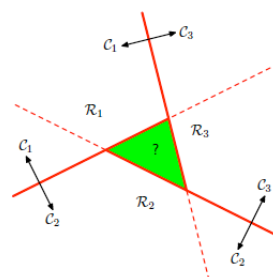
- Define $K - 1$ discrimination functions
- Function f_i ($1 \leq i \leq K - 1$) discriminates points belonging to class C_i from points belonging to all other classes: if $f_i(\mathbf{x}) > 0$ then $\mathbf{x} \in C_i$, otherwise $\mathbf{x} \notin C_i$
- The green region belongs to both \mathcal{R}_1 and \mathcal{R}_2



Linear discriminant functions in multiclass classification

Second approach

- Define $K(K - 1)/2$ discrimination functions, one for each pair of classes
- Function f_{ij} ($1 \leq i < j \leq K$) discriminates points which might belong to C_i from points which might belong to C_j
- Item \mathbf{x} is classified on a majority basis
- The green region is unassigned



Linear discriminant functions in multiclass classification

Third approach

- Define K linear functions

$$y_i(\mathbf{x}) = \mathbf{w}_i^T \mathbf{x} + w_{i0} \quad 1 \leq i \leq K$$

Item \mathbf{x} is assigned to class C_k iff $y_k(\mathbf{x}) > y_j(\mathbf{x})$ for all $j \neq k$: that is,

$$k = \underset{j}{\operatorname{argmax}} y_j(\mathbf{x})$$

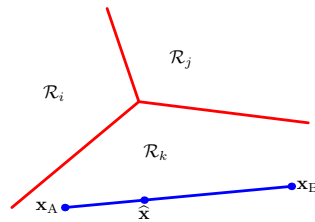
- Decision boundary between C_i and C_j : all points \mathbf{x} s.t. $y_i(\mathbf{x}) = y_j(\mathbf{x})$, a $D - 1$ -dimensional hyperplane

$$(\mathbf{w}_i - \mathbf{w}_j)^T \mathbf{x} + (w_{i0} - w_{j0}) = 0$$

Linear discriminant functions in multiclass classification

The resulting decision regions are connected and convex

- Given $\mathbf{x}_A, \mathbf{x}_B \in \mathcal{R}_k$ then $y_k(\mathbf{x}_A) > y_j(\mathbf{x}_A)$ and $y_k(\mathbf{x}_B) > y_j(\mathbf{x}_B)$, for all $j \neq k$
- Let $\hat{\mathbf{x}} = \lambda \mathbf{x}_A + (1 - \lambda) \mathbf{x}_B$, $0 \leq \lambda \leq 1$
- For all i , since y_i is linear for all, $y_i(\hat{\mathbf{x}}) = \lambda y_i(\mathbf{x}_A) + (1 - \lambda) y_i(\mathbf{x}_B)$
- Then, $y_k(\hat{\mathbf{x}}) > y_j(\hat{\mathbf{x}})$ for all $j \neq k$; that is, $\hat{\mathbf{x}} \in \mathcal{R}_k$



Generalized discriminant functions

- The definition can be extended to include terms relative to products of pairs of feature values (*Quadratic discriminant functions*)

$$y(\mathbf{x}) = w_0 + \sum_{i=1}^D w_i x_i + \sum_{i=1}^D \sum_{j=1}^i w_{ij} x_i x_j$$

$\frac{d(d+1)}{2}$ additional parameters wrt the $d + 1$ original ones: decision boundaries can be more complex

- In general, *generalized discrimination functions* through set of functions ϕ_1, \dots, ϕ_m

$$y(\mathbf{x}) = w_0 + \sum_{i=1}^M w_i \phi_i(\mathbf{x})$$

2 Least squares and classification

Linear discriminant functions and regression

- Assume classification with K classes
- Classes are represented through a 1-of- K coding scheme: set of variables z_1, \dots, z_K , class C_i coded by values $z_i = 1, z_k = 0$ for $k \neq i$
- Discriminant functions y_i are derived as linear regression functions with variables z_i as targets
- To each variable z_i a discriminant function $y_i(\mathbf{x}) = \mathbf{w}_i^T \mathbf{x} + w_{i0}$ is associated: \mathbf{x} is assigned to the class C_k s.t.

$$k = \underset{i}{\operatorname{argmax}} y_i(\mathbf{x})$$

- Then, $z_k(\mathbf{x}) = 1$ and $z_j(\mathbf{x}) = 0$ ($j \neq k$) if $k = \underset{i}{\operatorname{argmax}} y_i(\mathbf{x})$
- Group all parameters together as

$$\mathbf{y}(\mathbf{x}) = \mathbf{W}^T \bar{\mathbf{x}}$$

Linear discriminant functions and regression

- In general, a regression function provides an estimation of the target given the input $E[t|\mathbf{x}]$
- Value $y_i(\mathbf{x})$ can then be seen as an estimation of the conditional expectation $E[z_i|\mathbf{x}]$ of binary variable z_i given \mathbf{x}
- If we assume z_i is distributed according to a Bernoulli distribution, the expectation corresponds to the posterior probability

$$\begin{aligned} y_i(\mathbf{x}) &\simeq E[z_i|\mathbf{x}] \\ &= P(z_i = 1|\mathbf{x}) \cdot 1 + P(z_i = 0|\mathbf{x}) \cdot 0 \\ &= P(z_i = 1|\mathbf{x}) \\ &= P(C_i|\mathbf{x}) \end{aligned}$$

- However, $y_i(\mathbf{x})$ is not a probability itself (we may not assume it takes value only in the interval $[0, 1]$)

Learning functions y_i

- Given a training set \mathbf{X}, \mathbf{t} , a regression function can be derived by least squares
- An item in the training set is a pair $(\mathbf{x}_i, \mathbf{t}_i)$, $\mathbf{x}_i \in \mathbf{R}^D$ e $\mathbf{t}_i \in \{0, 1\}^K$
- $\mathbf{W} \in \mathbf{R}^{(D+1) \times K}$ is the matrix of parameters of all functions y_i : the i -th column represents the $D + 1$ parameters w_{i0}, \dots, w_{iD} of y_i

$$\bar{\mathbf{W}} = \begin{pmatrix} w_{10} & w_{20} & \cdots & w_{K0} \\ w_{11} & w_{21} & \cdots & w_{K1} \\ \vdots & \vdots & \ddots & \vdots \\ w_{1D} & w_{2D} & \cdots & w_{KD} \end{pmatrix}$$

- $\mathbf{y}(\mathbf{x}) = \mathbf{W}^T \bar{\mathbf{x}}$ with $\bar{\mathbf{x}} = (1, x_1, \dots, x_d)$

Learning functions y_i

- $\bar{\mathbf{X}} \in \mathbb{R}^{n \times (D+1)}$ is the matrix of feature values for all items in the training set

$$\bar{\mathbf{X}} = \begin{pmatrix} 1 & x_1^{(1)} & \dots & x_1^{(D)} \\ 1 & x_2^{(1)} & \dots & x_2^{(D)} \\ \vdots & \vdots & \ddots & \vdots \\ 1 & x_n^{(1)} & \dots & x_n^{(D)} \end{pmatrix}$$

- Then, for matrix $\bar{\mathbf{X}}\mathbf{W}$, of size $n \times K$, we have

$$(\bar{\mathbf{X}}\mathbf{W})_{ij} = w_{j0} + \sum_{k=1}^D x_i^{(k)} w_{jk} = y_j(\mathbf{x}_i)$$

Learning functions y_i

- $y_j(\mathbf{x}_i)$ is compared to item \mathbf{T}_{ij} in the matrix \mathbf{T} , of size $n \times K$, of target values, where row i is the 1-of- K coding of the class of item \mathbf{x}_i

$$(\bar{\mathbf{X}}\mathbf{W} - \mathbf{T})_{ik} = y_k(\mathbf{x}_i) - t_{ik}$$

- Let us consider the diagonal items of $(\bar{\mathbf{X}}\mathbf{W} - \mathbf{T})^T(\bar{\mathbf{X}}\mathbf{W} - \mathbf{T})$. Then,

$$((\bar{\mathbf{X}}\mathbf{W} - \mathbf{T})^T(\bar{\mathbf{X}}\mathbf{W} - \mathbf{T}))_{kk} = \sum_{i=1}^n (y_k(\mathbf{x}_i) - t_{ik})^2$$

That is,

$$((\bar{\mathbf{X}}\mathbf{W} - \mathbf{T})^T(\bar{\mathbf{X}}\mathbf{W} - \mathbf{T}))_{kk} = \sum_{\mathbf{x}_i \in C_k} (y_k(\mathbf{x}_i) - 1)^2 + \sum_{\mathbf{x}_i \notin C_k} y_k(\mathbf{x}_i)^2$$

Learning functions y_i

- Summing all elements on the diagonal of $(\bar{\mathbf{X}}\mathbf{W} - \mathbf{T})^T(\bar{\mathbf{X}}\mathbf{W} - \mathbf{T})$ provides the overall sum, on all items in the training set, of the squared differences between observed values and values computed by the model, with parameters \mathbf{W} , that is

$$\sum_{k=1}^K \sum_{i=1}^n (y_k(\mathbf{x}_i) - t_{ik})^2$$

- This corresponds to the *trace* of $(\bar{\mathbf{X}}\mathbf{W} - \mathbf{T})^T(\bar{\mathbf{X}}\mathbf{W} - \mathbf{T})$. Hence, we have to minimize:

$$E(\mathbf{W}) = \frac{1}{2} \text{tr}((\bar{\mathbf{X}}\mathbf{W} - \mathbf{T})^T(\bar{\mathbf{X}}\mathbf{W} - \mathbf{T}))$$

- Standard approach, solve

$$\frac{\partial E(\mathbf{W})}{\partial \mathbf{W}} = \mathbf{0}$$

Learning functions y_i

- It is possible to show that

$$\frac{\partial E(\mathbf{W})}{\partial \mathbf{W}} = \bar{\mathbf{X}}^T \bar{\mathbf{X}} \mathbf{W} - \bar{\mathbf{X}}^T \mathbf{T}$$

- From $\bar{\mathbf{X}}^T \bar{\mathbf{X}} \mathbf{W} - \bar{\mathbf{X}}^T \mathbf{T} = \mathbf{0}$ it results


$$\mathbf{W} = (\bar{\mathbf{X}}^T \bar{\mathbf{X}})^{-1} \bar{\mathbf{X}}^T \mathbf{T}$$

- and the set of discriminant functions


$$\mathbf{y}(\mathbf{x}) = \mathbf{W}^T \bar{\mathbf{x}} = \mathbf{T}^T \bar{\mathbf{X}} (\bar{\mathbf{X}}^T \bar{\mathbf{X}})^{-1} \bar{\mathbf{x}}$$

Some considerations

- Simple learning: closed form
- quite prone to outliers (magenta, this approach; green, logistic regression)



- poor precision for $K > 2$ (left, this approach; right, logistic regression)



3 Fisher' linear discriminant

Fisher' linear discriminant

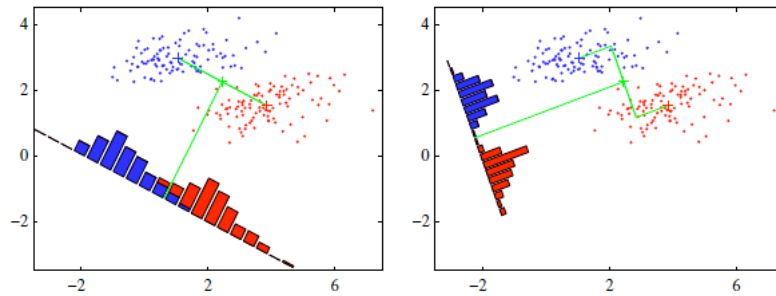
- The idea of *Linear Discriminant Analysis (LDA)* is to find a linear projection of the training set into a suitable subspace where classes are as linearly separated as possible
- A common approach is provided by *Fisher linear discriminant*, where all items in the training set (points in a D -dimensional space) are projected to one dimension, by means of a linear transformation of the type

$$y = \mathbf{w} \cdot \mathbf{x} = \mathbf{w}^T \mathbf{x}$$

where \mathbf{w} is the D -dimensional vector corresponding to the direction of projection (in the following, we will consider the one with unit norm).

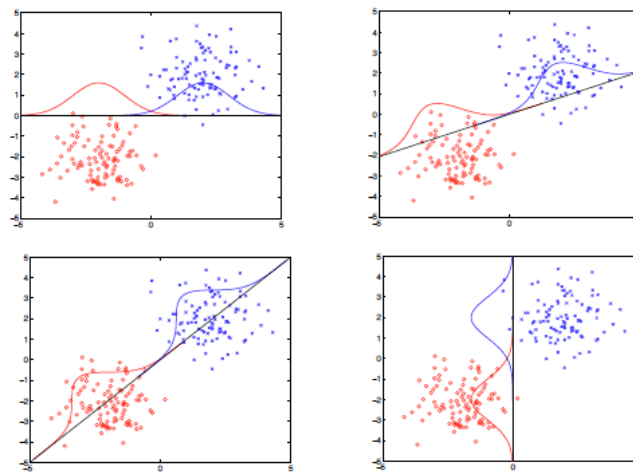
LDA

If $K = 2$, given a threshold \tilde{y} , item \mathbf{x} is assigned to C_1 iff its projection $y = \mathbf{w}^T \mathbf{x}$ is such that $y > \tilde{y}$; otherwise, \mathbf{x} is assigned to C_2 .



LDA

Different line directions, that is different parameters \mathbf{w} , may induce quite different separability properties.



Deriving \mathbf{w} in the binary case

Let n_1 be the number of items in the training set belonging to class C_1 and n_2 the number of items in class C_2 . The mean points of both classes are

$$\mathbf{m}_1 = \frac{1}{n_1} \sum_{\mathbf{x} \in C_1} \mathbf{x} \quad \mathbf{m}_2 = \frac{1}{n_2} \sum_{\mathbf{x} \in C_2} \mathbf{x}$$

A simple measure of the separation of classes, when the training set is projected onto a line, is the difference between the projections of their mean points

$$m_2 - m_1 = \mathbf{w}^T (\mathbf{m}_2 - \mathbf{m}_1)$$

where $m_i = \mathbf{w}^T \mathbf{m}_i$ is the projection of \mathbf{m}_i onto the line.

Deriving \mathbf{w} in the binary case

- We wish to find a line direction \mathbf{w} such that $m_2 - m_1$ is maximum
- $\mathbf{w}^T (\mathbf{m}_2 - \mathbf{m}_1)$ can be made arbitrarily large by multiplying \mathbf{w} by a suitable constant, at the same time maintaining the direction unchanged. To avoid this drawback, we consider unit vectors, introducing the constraint $\|\mathbf{w}\|_2 = \mathbf{w}^T \mathbf{w} = 1$

- This results into the constrained optimization problem

$$\max_{\mathbf{w}} \mathbf{w}^T (\mathbf{m}_2 - \mathbf{m}_1)$$

where $\mathbf{w}^T \mathbf{w} = 1$

- This can be transformed into an equivalent unconstrained optimization problem by means of *lagrangian multipliers*

$$\max_{\mathbf{w}, \lambda} \mathbf{w}^T (\mathbf{m}_2 - \mathbf{m}_1) + \lambda(1 - \mathbf{w}^T \mathbf{w})$$

Deriving \mathbf{w} in the binary case

Setting the gradient of the function wrt \mathbf{w} to $\mathbf{0}$

$$\frac{\partial}{\partial \mathbf{w}} (\mathbf{w}^T (\mathbf{m}_2 - \mathbf{m}_1) + \lambda(1 - \mathbf{w}^T \mathbf{w})) = \mathbf{m}_2 - \mathbf{m}_1 + 2\lambda \mathbf{w} = \mathbf{0}$$

results into

$$\mathbf{w} = \frac{\mathbf{m}_2 - \mathbf{m}_1}{2\lambda}$$

Deriving \mathbf{w} in the binary case

Setting the derivative wrt λ to 0

$$\frac{\partial}{\partial \lambda} (\mathbf{w}^T (\mathbf{m}_2 - \mathbf{m}_1) + \lambda(1 - \mathbf{w}^T \mathbf{w})) = 1 - \mathbf{w}^T \mathbf{w} = 0$$

results into

$$1 - \mathbf{w}^T \mathbf{w} = 1 - \frac{(\mathbf{m}_2 - \mathbf{m}_1)^T (\mathbf{m}_2 - \mathbf{m}_1)}{4\lambda^2} = 0$$

that is

$$\lambda = \frac{\sqrt{(\mathbf{m}_2 - \mathbf{m}_1)^T (\mathbf{m}_2 - \mathbf{m}_1)}}{2} = \frac{\|\mathbf{m}_2 - \mathbf{m}_1\|_2}{2}$$

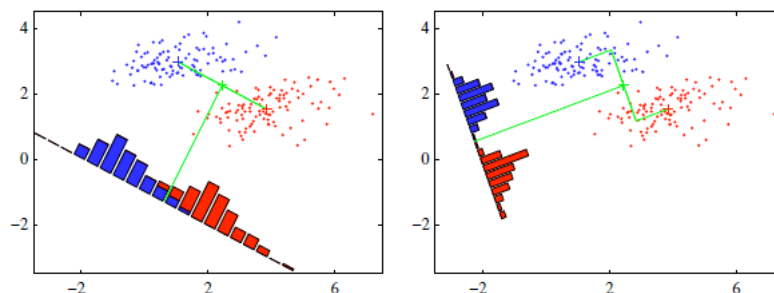
Combining with the result for the gradient, we get

$$\mathbf{w} = \frac{\mathbf{m}_2 - \mathbf{m}_1}{\|\mathbf{m}_2 - \mathbf{m}_1\|_2}$$

Deriving \mathbf{w} in the binary case

The best direction \mathbf{w} of the line, wrt the measure considered, is the one from \mathbf{m}_1 to \mathbf{m}_2 .

However, this may result in a poor separation of classes.



Projections of classes are dispersed (high variance) along the direction of $\mathbf{m}_1 - \mathbf{m}_2$. This may result in a large overlap.

Deriving \mathbf{w} in the binary case: refinement

- Choose directions s.t. classes projections show as little dispersion as possible
- Possible in the case that the amount of class dispersion changes wrt different directions, that is if the distribution of points in the class is elongated
- We wish then to maximize a function which:
 - is growing wrt the separation between the projected classes (for example, their mean points)
 - is decreasing wrt the dispersion of the projections of points of each class

Deriving \mathbf{w} in the binary case: refinement

- The *within-class variance* of the projection of class C_i ($i = 1, 2$) is defined as

$$s_i^2 = \sum_{\mathbf{x} \in C_i} (\mathbf{w}^T \mathbf{x} - m_i)^2$$

The total within-class variance is defined as $s_1^2 + s_2^2$

- Given a direction \mathbf{w} , the *Fisher criterion* is the ratio between the (squared) class separation and the overall within-class variance, along that direction

$$J(\mathbf{w}) = \frac{(m_2 - m_1)^2}{s_1^2 + s_2^2}$$

- Indeed, $J(\mathbf{w})$ grows wrt class separation and decreases wrt within-class variance

Deriving \mathbf{w} in the binary case: refinement

Let $\mathbf{S}_1, \mathbf{S}_2$ be the *within-class covariance matrices*, defined as

$$\mathbf{S}_i = \sum_{\mathbf{x} \in C_i} (\mathbf{x} - \mathbf{m}_i)(\mathbf{x} - \mathbf{m}_i)^T$$

Then,

$$\begin{aligned} s_i^2 &= \sum_{\mathbf{x} \in C_i} (\mathbf{w}^T \mathbf{x} - m_i)^2 = \sum_{\mathbf{x} \in C_i} (\mathbf{w}^T \mathbf{x} - \mathbf{w}^T \mathbf{m}_i)^2 \\ &= \sum_{\mathbf{x} \in C_i} (\mathbf{w}^T \mathbf{x} - \mathbf{w}^T \mathbf{m}_i)(\mathbf{w}^T \mathbf{x} - \mathbf{w}^T \mathbf{m}_i) \\ &= \sum_{\mathbf{x} \in C_i} (\mathbf{w}^T \mathbf{x} - \mathbf{w}^T \mathbf{m}_i)(\mathbf{x}^T \mathbf{w} - \mathbf{m}_i^T \mathbf{w}) \\ &= \sum_{\mathbf{x} \in C_i} (\mathbf{w}^T (\mathbf{x} - \mathbf{m}_i)) ((\mathbf{x} - \mathbf{m}_i)^T \mathbf{w}) \\ &= \sum_{\mathbf{x} \in C_i} \mathbf{w}^T (\mathbf{x} - \mathbf{m}_i)(\mathbf{x} - \mathbf{m}_i)^T \mathbf{w} \\ &= \mathbf{w}^T \left(\sum_{\mathbf{x} \in C_i} (\mathbf{x} - \mathbf{m}_i)(\mathbf{x} - \mathbf{m}_i)^T \right) \mathbf{w} = \mathbf{w}^T \mathbf{S}_i \mathbf{w} \end{aligned}$$

Deriving \mathbf{w} in the binary case: refinement

Let also $\mathbf{S}_W = \mathbf{S}_1 + \mathbf{S}_2$ be the *total within-class covariance matrix* and

$$\mathbf{S}_B = (\mathbf{m}_2 - \mathbf{m}_1)(\mathbf{m}_2 - \mathbf{m}_1)^T$$

be the *between-class covariance matrix*.

Then,

$$\begin{aligned}
 J(\mathbf{w}) &= \frac{(m_2 - m_1)^2}{s_1^2 + s_2^2} = \frac{(\mathbf{w}^T \mathbf{m}_2 - \mathbf{w}^T \mathbf{m}_1)^2}{\mathbf{w}^T \mathbf{S}_1 \mathbf{w} + \mathbf{w}^T \mathbf{S}_2 \mathbf{w}} \\
 &= \frac{(\mathbf{w}^T \mathbf{m}_2 - \mathbf{w}^T \mathbf{m}_1)(\mathbf{w}^T \mathbf{m}_2 - \mathbf{w}^T \mathbf{m}_1)}{\mathbf{w}^T \mathbf{S}_1 \mathbf{w} + \mathbf{w}^T \mathbf{S}_2 \mathbf{w}} \\
 &= \frac{\mathbf{w}^T (\mathbf{m}_2 - \mathbf{m}_1)(\mathbf{m}_2 - \mathbf{m}_1)^T \mathbf{w}}{\mathbf{w}^T \mathbf{S}_1 \mathbf{w} + \mathbf{w}^T \mathbf{S}_2 \mathbf{w}} \\
 &= \frac{\mathbf{w}^T \mathbf{S}_B \mathbf{w}}{\mathbf{w}^T \mathbf{S}_W \mathbf{w}}
 \end{aligned}$$

Deriving \mathbf{w} in the binary case: refinement

As usual, $J(\mathbf{w})$ is maximized wrt \mathbf{w} by setting its gradient to $\mathbf{0}$

$$\frac{\partial}{\partial \mathbf{w}} \frac{\mathbf{w}^T \mathbf{S}_B \mathbf{w}}{\mathbf{w}^T \mathbf{S}_W \mathbf{w}} = 2 \frac{(\mathbf{w}^T \mathbf{S}_B \mathbf{w}) \mathbf{S}_W \mathbf{w} - (\mathbf{w}^T \mathbf{S}_W \mathbf{w}) \mathbf{S}_B \mathbf{w}}{(\mathbf{w}^T \mathbf{S}_W \mathbf{w})(\mathbf{w}^T \mathbf{S}_W \mathbf{w})^T}$$

which results into

$$(\mathbf{w}^T \mathbf{S}_B \mathbf{w}) \mathbf{S}_W \mathbf{w} = (\mathbf{w}^T \mathbf{S}_W \mathbf{w}) \mathbf{S}_B \mathbf{w}$$

Deriving \mathbf{w} in the binary case: refinement

Observe that:

- $\mathbf{w}^T \mathbf{S}_B \mathbf{w}$ is a scalar, say c_B
- $\mathbf{w}^T \mathbf{S}_W \mathbf{w}$ is a scalar, say c_W
- $(\mathbf{m}_2 - \mathbf{m}_1)^T \mathbf{w}$ is a scalar, say c_m

Then, the condition $(\mathbf{w}^T \mathbf{S}_B \mathbf{w}) \mathbf{S}_W \mathbf{w} = (\mathbf{w}^T \mathbf{S}_W \mathbf{w}) \mathbf{S}_B \mathbf{w}$ can be written as

$$c_B \mathbf{S}_W \mathbf{w} = c_W \mathbf{S}_B \mathbf{w} = c_W (\mathbf{m}_2 - \mathbf{m}_1)(\mathbf{m}_2 - \mathbf{m}_1)^T \mathbf{w} = c_W (\mathbf{m}_2 - \mathbf{m}_1) c_m$$

which results into

$$\mathbf{w} = \frac{c_W c_m}{c_B} \mathbf{S}_W^{-1} (\mathbf{m}_2 - \mathbf{m}_1)$$

Since we are interested into the direction of \mathbf{w} , that is in any vector proportional to \mathbf{w} , we may consider the solution

$$\hat{\mathbf{w}} = \mathbf{S}_W^{-1} (\mathbf{m}_2 - \mathbf{m}_1) = (\mathbf{S}_1 + \mathbf{S}_2)^{-1} (\mathbf{m}_2 - \mathbf{m}_1)$$

Deriving \mathbf{w} in the binary case: choosing a threshold

Possible approach:

- model $p(y|C_i)$ as a gaussian: derive mean and variance by maximum likelihood

$$m_i = \frac{1}{n_i} \sum_{\mathbf{x} \in C_i} w^T \mathbf{x} \quad \sigma_i^2 = \frac{1}{n_i - 1} \sum_{\mathbf{x} \in C_i} (w^T \mathbf{x} - m_i)^2$$

where n_i is the number of items in training set belonging to class C_i

- derive the class probabilities

$$p(C_i|y) \propto p(y|C_i)p(C_i) = p(y|C_i) \frac{n_i}{n_1 + n_2} \propto n_i e^{-\frac{(y-m_i)^2}{2\sigma_i^2}}$$

- the threshold \tilde{y} can be derived as the minimum y such that

$$\frac{p(C_2|y)}{p(C_1|y)} = \frac{n_2 p(y|C_2)}{n_1 p(y|C_1)} > 1$$

Perceptron

- Introduced in the '60s, at the basis of the neural network approach
- Simple model of a single neuron
- Hard to evaluate in terms of probability
- Works only in the case that classes are linearly separable

Definition

It corresponds to a binary classification model where an item \mathbf{x} is first transformed by a non linear function ϕ and the classified on the basis of the sign of the obtained value. That is,

$$y(\mathbf{x}) = f(\mathbf{w}^T \phi(\mathbf{x}))$$

$f()$ is essentially the sign function

$$f(i) = \begin{cases} -1 & \text{if } i < 0 \\ 1 & \text{if } i \geq 0 \end{cases}$$

The resulting model is a particular generalized linear model. A special case is the one when ϕ is the identity, that is $y(\mathbf{x}) = f(\mathbf{w}^T \mathbf{x})$.

By the definition of the model, $y(\mathbf{x})$ can only be ± 1 : we denote $y(\mathbf{x}) = 1$ as $\mathbf{x} \in C_1$ and $y(\mathbf{x}) = -1$ as $\mathbf{x} \in C_2$.

To each element \mathbf{x}_i in the training set, a target value is then associated $t_i \in \{-1, 1\}$.

Cost function

- A natural definition of the cost function would be the number of misclassified elements in the training set
- This would result into a piecewise constant function and gradient optimization could not be applied (we would have zero gradient almost everywhere)
- A better choice is using a piecewise linear function as cost function

Cost function

We would like to find a vector of parameters \mathbf{w} such that, for any \mathbf{x}_i , $\mathbf{w}^T \mathbf{x}_i > 0$ if $\mathbf{x}_i \in C_1$ and $\mathbf{w}^T \mathbf{x}_i < 0$ if $\mathbf{x}_i \in C_2$: in short, $\mathbf{w}^T \mathbf{x}_i t_i > 0$.

Each element \mathbf{x}_i provides a contribution to the cost function as follows

1. 0 if \mathbf{x}_i is classified correctly by the model
2. $-\mathbf{w}^T \mathbf{x}_i t_i > 0$ if \mathbf{x}_i is misclassified

Let \mathcal{M} be the set of misclassified elements. Then the cost is

$$E_p(\mathbf{w}) = - \sum_{\mathbf{x}_i \in \mathcal{M}} \mathbf{w}^T \phi(\mathbf{x}_i) t_i$$

The contribution of \mathbf{x}_i to the cost is 0 if $\mathbf{x}_i \notin \mathcal{M}$ and it is a linear function of \mathbf{w} otherwise

Gradient optimization

The minimum of $E_p(\mathbf{w})$ can be found through gradient descent

$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} - \eta \frac{\partial E_p(\mathbf{w})}{\partial \mathbf{w}} \Big|_{\mathbf{w}^{(k)}}$$

the gradient of the cost function wrt to \mathbf{w} is

$$\frac{\partial E_p(\mathbf{w})}{\partial \mathbf{w}} = - \sum_{\mathbf{x}_i \in \mathcal{M}} \phi(\mathbf{x}_i) t_i$$

Then gradient descent can be expressed as

$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} + \eta \sum_{\mathbf{x}_i \in \mathcal{M}_k} \phi(\mathbf{x}_i) t_i$$

where \mathcal{M}_k denotes the set of points misclassified by the model with parameter $\mathbf{w}^{(k)}$

Gradient optimization

Online (or stochastic gradient descent): at each step, only the gradient wrt a single item is considered

$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} + \eta \phi(\mathbf{x}_i) t_i$$

where $\mathbf{x}_i \in \mathcal{M}_k$ and the *scale factor* $\eta > 0$ controls the impact of a badly classified item on the cost function

The method works by circularly iterating on all elements and applying the above formula.

Gradient optimization

Initialize \mathbf{w}^0

$k := 0$ **repeat**

$k := k + 1$

$i := (k \bmod n) + 1$

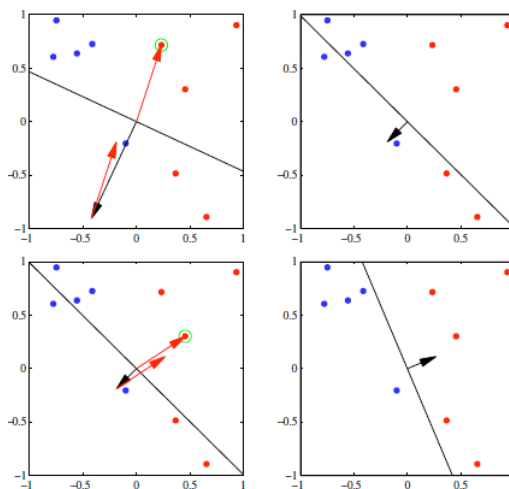
$y := f(\mathbf{w}^T \phi(\mathbf{x}_i)) t_i$

if $y > 0$ **then** $\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)}$

else $\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} + \eta \phi(\mathbf{x}_i) t_i$

until all elements are well classified

Gradient optimization



In black, decision boundary and corresponding parameter vector \mathbf{w} ; in red misclassified item vector $\phi(\mathbf{x}_i)$, added by the algorithm to the parameter vector as $\eta\phi(\mathbf{x}_i)$

Gradient optimization

At each step, if \mathbf{x}_i is well classified then $\mathbf{w}^{(k)}$ is unchanged; else, its contribution to the cost is modified as follows

$$\begin{aligned} -(\mathbf{w}^{(k+1)})^T \phi(\mathbf{x}_i) t_i &= -(\mathbf{w}^{(k)})^T \phi(\mathbf{x}_i) t_i - \eta(\phi(\mathbf{x}_i) t_i)^T \phi(\mathbf{x}_i) t_i \\ &= -(\mathbf{w}^{(k)})^T \phi(\mathbf{x}_i) t_i - \eta \|\phi(\mathbf{x}_i)\|^2 \\ &< -(\mathbf{w}^{(k)})^T \phi(\mathbf{x}_i) t_i \end{aligned}$$

This contribution is decreasing, however this does not guarantee the convergence of the method, since the cost function could increase due to some other element becoming misclassified if $\mathbf{w}^{(k+1)}$ is used

Perceptron convergence theorem

It is possible to prove that, in the case the classes are linearly separable, the algorithm converges to the correct solution in a finite number of steps.

Let $\hat{\mathbf{w}}$ be a solution (that is, it discriminates C_1 and C_2): if \mathbf{x}_{k+1} is the element considered at iteration $(k+1)$ and it is misclassified, then

$$\mathbf{w}^{(k+1)} - \alpha \hat{\mathbf{w}} = (\mathbf{w}^{(k)} - \alpha \hat{\mathbf{w}}) + \eta \phi(\mathbf{x}_{k+1}) t_{k+1}$$

where $\alpha > 0$ is a constant, to be specified later

Perceptron convergence theorem

By squaring left and right expressions of the above formula, we get

$$\begin{aligned} \|\mathbf{w}^{(k+1)} - \alpha \hat{\mathbf{w}}\|^2 &= \|\mathbf{w}^{(k)} - \alpha \hat{\mathbf{w}}\|^2 + \eta^2 \|\phi(\mathbf{x}_{k+1})\|^2 + 2\eta(\mathbf{w}^{(k)} - \alpha \hat{\mathbf{w}})^T \phi(\mathbf{x}_{k+1}) t_{k+1} = \\ &= \|\mathbf{w}^{(k)} - \alpha \hat{\mathbf{w}}\|^2 + \eta^2 \|\phi(\mathbf{x}_{k+1})\|^2 + 2\eta(\mathbf{w}^{(k)})^T \phi(\mathbf{x}_{k+1}) t_{k+1} - 2\eta\alpha \hat{\mathbf{w}}^T \phi(\mathbf{x}_{k+1}) t_{k+1} \end{aligned}$$

Since \mathbf{x}_{k+1} was misclassified by hypothesis, $(\mathbf{w}^{(k)})^T \phi(\mathbf{x}_{k+1}) t_{k+1} < 0$ and

$$\|\mathbf{w}^{(k+1)} - \alpha \hat{\mathbf{w}}\|^2 < \|\mathbf{w}^{(k)} - \alpha \hat{\mathbf{w}}\|^2 + \eta^2 \|\phi(\mathbf{x}_{k+1})\|^2 - 2\eta\alpha \hat{\mathbf{w}}^T \phi(\mathbf{x}_{k+1}) t_{k+1}$$

Perceptron convergence theorem

Let γ be the minimum value of the signed dot product of $\hat{\mathbf{w}}$ with $\phi(\mathbf{x}_i)$ for some element \mathbf{x}_i , where the sign depends on the class of \mathbf{x}_i

$$\gamma = \min_i (\hat{\mathbf{w}}^T \phi(\mathbf{x}_i) t_i) = \min_i |\hat{\mathbf{w}}^T \phi(\mathbf{x}_i)| > 0$$

Let δ be the length of the longest $\phi(\mathbf{x}_i)$

$$\delta^2 = \max_i \|\phi(\mathbf{x}_i)\|^2$$

Then,

$$\|\mathbf{w}^{(k+1)} - \alpha \hat{\mathbf{w}}\|^2 < \|\mathbf{w}^{(k)} - \alpha \hat{\mathbf{w}}\|^2 + \eta^2 \delta^2 - 2\eta\alpha\gamma$$

Perceptron convergence theorem

By setting

$$\alpha = \frac{\eta \delta^2}{\gamma}$$

we get

$$\left\| \mathbf{w}^{(k+1)} - \alpha \hat{\mathbf{w}} \right\|^2 < \left\| \mathbf{w}^{(k)} - \alpha \hat{\mathbf{w}} \right\|^2 - \eta^2 \delta^2$$

As can be seen, the squared distance between $\mathbf{w}^{(k+1)}$ and $\hat{\mathbf{w}}$ decreases at each step of an amount greater than $\eta^2 \delta^2$

Perceptron convergence theorem

Iterating the above properties on all steps,

$$\left\| \mathbf{w}^{(k+1)} - \alpha \hat{\mathbf{w}} \right\|^2 < \left\| \mathbf{w}^{(0)} - \alpha \hat{\mathbf{w}} \right\|^2 - (k+1)\eta^2 \delta^2$$

Note that, after

$$\bar{k} = \frac{\left\| \mathbf{w}^{(0)} - \alpha \hat{\mathbf{w}} \right\|^2}{\eta^2 \delta^2} - 1$$

steps we get

$$\left\| \mathbf{w}^{(0)} - \alpha \hat{\mathbf{w}} \right\|^2 - (k+1)\eta^2 \delta^2 = 0$$

So, after at most \bar{k} updates of \mathbf{w} , a decision boundary has been derived

Perceptron convergence theorem

Setting $\mathbf{w}^{(0)} = \mathbf{0}$, we have

$$\bar{k} = \frac{\alpha^2}{\eta^2 \delta^2} \|\hat{\mathbf{w}}\|^2 - 1 = \frac{\delta^2}{\gamma^2} \|\hat{\mathbf{w}}\|^2 - 1 = \frac{\max_i \|\phi(\mathbf{x}_i)\|^2}{(\min_i (\hat{\mathbf{w}}^T \phi(\mathbf{x}_i)))^2} \|\hat{\mathbf{w}}\|^2 - 1$$

The number of required step is large if $\min_i (\hat{\mathbf{w}}^T \phi(\mathbf{x}_i))$ is small, that is if there exists some \mathbf{x}_i such that $\phi(\mathbf{x}_i)$ is (almost) orthogonal to $\hat{\mathbf{w}}$.